

Assessing and Benchmarking zkVMs: Insights into Performance and Scalability

Lawrence Lim, Nihar Shah

2024-05-13

Presentation Structure

- ▶ Introduction
 - ▶ What is a zkVM?
 - ▶ Table 1 - zkVM Architecture
 - ▶ The zkVM Landscape
 - ▶ Why RISC-V?
- ▶ zkVM Technical Overview
 - ▶ Frontends and Backends of zkVMs
 - ▶ Why do we care?
 - ▶ How to optimize zkVM performance
- ▶ zkVM Comparison
 - ▶ Table 2 - Component Comparison
 - ▶ Design Tradeoff Comparison
 - ▶ Benchmarking Rationale
 - ▶ Benchmarking Results
- ▶ Conclusion
- ▶ References

Introduction

Our project is an analysis and evaluation of zkVM construction and performance, benchmarking how the performance of different zkVMs scales with the memory usage of applications.

What is a zkVM?

A zkVM, is simply a VM implemented as a circuit for a zero-knowledge proof (zkp) system. So, instead of proving the execution of a program, as one would normally do in zkp systems, you prove the execution of the bytecode of a given Instruction Set Architecture (ISA).

There are a few types of zkVMs available on the market targeting different ISAs with various practical tradeoffs.

Table 1 - zkVM Architecture

	Existing Expertise / Tooling	Blockchain Focused	Performant
Mainstream ISAs RISC-V, WASM, MIPS	Lots	No	Maybe
EVM-Equivalent EVM Bytecode	Some	Yes	No
ZK-optimized New Instruction Set	No	Yes	Yes

The zkVM Landscape

EVM Equivalent:

- ▶ Type 1: Taiko
- ▶ Type 2-3: Scroll, Polygon zkEVM
- ▶ Type 4: zkSync

Mainstream ISAs

- ▶ RISC-V: Succinct's SP1, a16z's Jolt, RISC-0
- ▶ WASM: zkWASM
- ▶ MIPS: zkMIPS

ZK Optimized

- ▶ Polygon Miden, Starknet Cairo

Why RISC-V?

Extremely popular compile target for many programming languages (Rust, C++, LLVM). Open sourced. RISC vs. CISC → RISC has less instructions and is therefore easier to arithmetize and prove than x86 assembly for example.

zkVM Technical Overview

Now let's take a closer look at the frontend and backend components of zkVMs.

Frontends and Backends of zkVMs

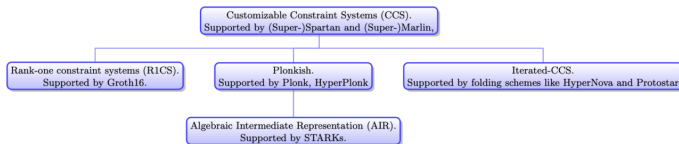


Figure 1: Intermediate representations (i.e., kinds of circuits) and the back-ends that can prove statements about them. The top-most box is the most general kind of circuit amongst those depicted, and each child is a special case of its parent. A subtlety not portrayed in the depicted taxonomy is that some SNARKs (such as Spartan and STARKs) can avoid having an honest party pre-process circuits if the circuits have repeated structure, while others cannot (Groth16, Marlin, Plonk).

Figure 1: Untitled

Frontend - Arithmetization Scheme

In general, arithmetization cannot be done manually except for elementary programs. Besides, the use of naïve arithmetization can lead to significant overhead. To deal with this, dedicated compilers accepting high-level programming languages have been developed.

Frontend - Precompiles

- ▶ Performance Issue: zkVMs operate significantly slower compared to running programs without proving overhead.
- ▶ Use of Precompiles: To improve efficiency, deployed zkVMs utilize “precompiles” — hand-optimized protocols for frequently used computations like hashing and signature verification.
- ▶ Risks of Overreliance: Relying heavily on precompiles can be problematic as designing these optimized protocols is the exact labor-intensive and error-prone process zkVMs aim to eliminate.

Backend - PCS & PIOP

The backend for proof system involves what we have learned in class, composed of two components: a PCS and a PIOP. Something notable about the interaction between the frontend and backend: most SNARKs can be easily tweaked to support both Plonkish and AIR with the exception of Groth16 which can only support R1CS.

Backend - Field Sizes

- ▶ Field Size Trade-offs: Operations in smaller fields are generally faster than in larger fields.
- ▶ Small vs. Large Fields: Using fields slightly smaller than 256 bits, like Goldilocks, can complicate operations with 256-bit numbers, requiring two field elements per value and roughly doubling prover costs.
- ▶ R1CS and Field Size: Currently, the R1CS system is constrained to larger fields, limiting flexibility in field size choice.

Backend - FRI Expansion Factor

The FRI blowup factor is a tunable parameter that allows you to adjust the cost to be more on the proving side or verifying side. A relatively low blowup factor leads to less prover time with larger proofs and a larger blowup factor leads to high cost to prove with smaller proof size.

Backend - Lookup Arguments

- ▶ Precomputed Outputs: This technique involves precomputing outputs for all possible inputs of a bitwise instruction.
- ▶ Efficient Verification: In a zkVM, a cost-effective SNARK operation, known as “the lookup,” verifies that the current instruction matches an entry in the precomputed table.
- ▶ Reduced Costs: Using lookup arguments decreases the cost of proving the instruction.

Why do we care?

It's important to distinguish between the backend and frontends of SNARKs to make clear assertions of the performance tradeoffs between various arithmetization schemes and SNARK backends. Failure to distinguish between them can result in misconceptions about performance and other characteristics of SNARKs

How to optimize zkVM performance

By efficient, we are almost always referring to proof generation time. Verifier time is about the same because we can use recursion to quickly verify proofs. Here are the options:

- ▶ Lookup tables.
- ▶ SNARK-friendly cryptographic primitives (such as Rescue, SAVER or Poseidon).
- ▶ Concurrent proof generation.
- ▶ Hardware acceleration (such as using GPU or FPGA).

zkVM Comparison

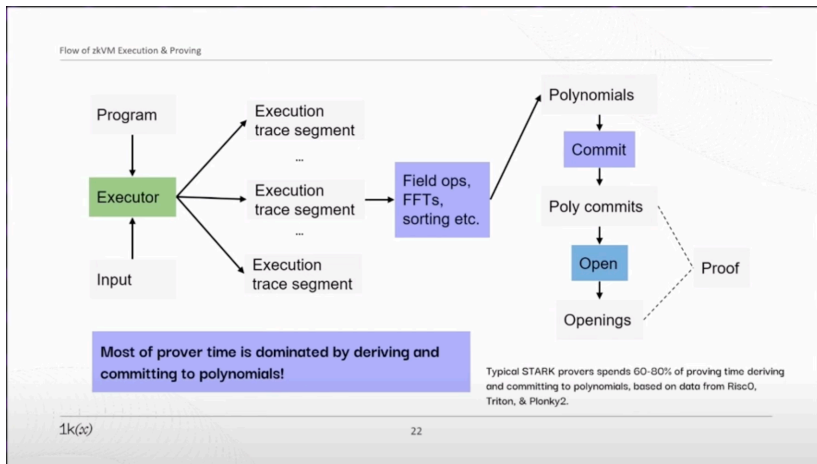


Figure 2: Untitled

Table 2 - Component Comparison

	RISC0	SP1	Jolt
PCS	FRI	FRI	Hyrax
Lookups	Plookup	Plookup?	Lasso
Field Size	~31-bit (baby bear)	~64-bit (goldilocks)?	~256-bit
Recursive Proofs	Yes	Yes	No
Precompiles	No?	Yes	No
Optimized for GPU	Yes	No	No
Arithmetization	AIR	AIR	R1CS
FRI Exp. Rate	4	2	N/A
SNARK Prover	Plonky2 STARK?	Plonky3 STARK	Spartan

Design Tradeoff Comparison

Benchmarking Rationale

...

Benchmarking Results

...

Conclusion

- ▶ Benchmarks are misleading
- ▶ Precompiles are complicated
- ▶ Naming is hard
- ▶ Recursion is helpful

References